


```
from sklearn.preprocessing import MinMaxScaler
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Input
from tensorflow.keras.layers import Dense
from tensorflow.keras.layers import LeakyReLU
from tensorflow.keras.layers import BatchNormalization
from tensorflow.keras.models import load_model
```


```
import numpy as np
import pandas as pd
train = pd.read_csv('UNSW_NB15_training-set.csv')
test = pd.read_csv('UNSW_NB15_testing-set.csv')
from sklearn.preprocessing import OrdinalEncoder
ord_enc = OrdinalEncoder()
```

```
train['proto_code'] = ord_enc.fit_transform(train[['proto']])
train[['proto','proto_code']].head(175341)
train['state_code'] = ord_enc.fit_transform(train[['state']])
train[['state','state_code']].head(175341)
train['service_code'] = ord_enc.fit_transform(train[['service']])
train[['service','service_code']].head(175341)
```

	service	service_code	
0	-	0.0	
1	-	0.0	
2	-	0.0	
3	ftp	3.0	
4	-	0.0	
...	
175336	dns	2.0	
175337	-	0.0	
175338	dns	2.0	
175339	dns	2.0	
175340	dns	2.0	

175341 rows × 2 columns

```
test['proto_code'] = ord_enc.fit_transform(test[['proto']])
test[['proto','proto_code']].head(82332)
test['state_code'] = ord_enc.fit_transform(test[['state']])
test[['state','state_code']].head(82332)
test['service_code'] = ord_enc.fit_transform(test[['service']])
test[['service','service_code']].head(82332)
```

	service	service_code	
0	-	0.0	
1	-	0.0	
2	-	0.0	
3	-	0.0	
4	-	0.0	
...	
82327	-	0.0	
82328	-	0.0	
82329	-	0.0	
82330	-	0.0	
82331	-	0.0	

82332 rows × 2 columns

```
x1 = train[['dur','spkts','dpkts','proto_code','state_code','service_code',
            'sbytes','dbytes','rate','sttl','dttl','sload','dload','sloss','dloss',
            'sinpkt','dinpkt','sjit','djit','swin','stcpb',
            'dtcpb','dwin','tcprtt','synack','ackdat','smean','dmean','trans_depth','response_body_len','ct_srv_src','ct_state_t
```

```
        , 'ct_dst_ltm', 'ct_src_dport_ltm', 'ct_dst_sport_ltm', 'ct_dst_src_ltm', 'is_ftp_login',
        'ct_ftp_cmd', 'ct_flw_http_mthd', 'ct_src_ltm', 'ct_srv_dst', 'is_sm_ips_ports']]
y1 = train['label']
x2 = test[['dur', 'spkts', 'dpkts', 'proto_code', 'state_code', 'service_code',
        'sbytes', 'dbytes', 'rate', 'sttl', 'dttl', 'sload', 'dload', 'sloss', 'dloss',
        'sinpkt', 'dinpkt', 'sjit', 'djit', 'swin', 'stcpb',
        'dtcpb', 'dwin', 'tcprrt', 'synack', 'ackdat', 'smean', 'dmean', 'trans_depth', 'response_body_len', 'ct_srv_src', 'ct_state_t
        , 'ct_dst_ltm', 'ct_src_dport_ltm', 'ct_dst_sport_ltm', 'ct_dst_src_ltm', 'is_ftp_login',
        'ct_ftp_cmd', 'ct_flw_http_mthd', 'ct_src_ltm', 'ct_srv_dst', 'is_sm_ips_ports']]
y2 = test['label']
```

```
model = MinMaxScaler()
model.fit(x1)
x1 = model.transform(x1)
x2 = model.transform(x2)
```

```
n_inputs = x1.shape[1]
n_inputs = x2.shape[1]
```

```
input_data_shape = Input(shape=(n_inputs,))
encoder = Dense(n_inputs*2)(input_data_shape)
encoder = BatchNormalization()(encoder)
encoder = LeakyReLU()(encoder)
encoder = Dense(n_inputs)(encoder)
encoder = BatchNormalization()(encoder)
encoder = LeakyReLU()(encoder)
n_bottleneck = round(float(n_inputs)/2.0)
bottleneck = Dense(n_bottleneck)(encoder)
decoder = Dense(n_inputs)(bottleneck)
decoder = BatchNormalization()(decoder)
decoder = LeakyReLU()(decoder)
decoder = Dense(n_inputs*2)(bottleneck)
decoder = BatchNormalization()(decoder)
decoder = LeakyReLU()(decoder)
```

```
output = Dense(n_inputs, activation = 'linear')(decoder)
model_AE = Model(inputs = input_data_shape, outputs = output)
model_AE.compile(optimizer = 'adam', loss='mse')
model_AE.summary()
```

Model: "model_4"

Layer (type)	Output Shape	Param #
=====		
input_3 (InputLayer)	[(None, 42)]	0
dense_12 (Dense)	(None, 84)	3612
batch_normalization_8 (Batch Normalization)	(None, 84)	336
leaky_re_lu_8 (LeakyReLU)	(None, 84)	0
dense_13 (Dense)	(None, 42)	3570
batch_normalization_9 (Batch Normalization)	(None, 42)	168
leaky_re_lu_9 (LeakyReLU)	(None, 42)	0
dense_14 (Dense)	(None, 21)	903
dense_16 (Dense)	(None, 84)	1848
batch_normalization_11 (Batch Normalization)	(None, 84)	336
leaky_re_lu_11 (LeakyReLU)	(None, 84)	0
dense_17 (Dense)	(None, 42)	3570
=====		
Total params: 14,343		
Trainable params: 13,923		
Non-trainable params: 420		

```
history = model_AE.fit(x1,x1,epochs = 50,batch_size = 16,verbose = 2,validation_data = (x2,x2))
10959/10959 - 30s - loss: 3.9789e-04 - val_loss: 6.7395e-04 - 30s/epoch - 3ms/step
Epoch 22/50
10959/10959 - 29s - loss: 3.9789e-04 - val_loss: 6.7395e-04 - 29s/epoch - 3ms/step
Epoch 22/50
```

```
Epoch 23/50
10959/10959 - 29s - loss: 3.9491e-04 - val_loss: 6.7184e-04 - 29s/epoch - 3ms/step
Epoch 24/50
10959/10959 - 30s - loss: 3.7757e-04 - val_loss: 6.2764e-04 - 30s/epoch - 3ms/step
Epoch 25/50
10959/10959 - 35s - loss: 3.8318e-04 - val_loss: 7.8785e-04 - 35s/epoch - 3ms/step
Epoch 26/50
10959/10959 - 29s - loss: 3.7458e-04 - val_loss: 6.7092e-04 - 29s/epoch - 3ms/step
Epoch 27/50
10959/10959 - 30s - loss: 3.6294e-04 - val_loss: 7.2804e-04 - 30s/epoch - 3ms/step
Epoch 28/50
10959/10959 - 34s - loss: 3.5356e-04 - val_loss: 6.1226e-04 - 34s/epoch - 3ms/step
Epoch 29/50
10959/10959 - 29s - loss: 3.6016e-04 - val_loss: 8.6791e-04 - 29s/epoch - 3ms/step
Epoch 30/50
10959/10959 - 35s - loss: 3.6324e-04 - val_loss: 7.0834e-04 - 35s/epoch - 3ms/step
Epoch 31/50
10959/10959 - 35s - loss: 3.4638e-04 - val_loss: 8.2770e-04 - 35s/epoch - 3ms/step
Epoch 32/50
10959/10959 - 30s - loss: 3.4158e-04 - val_loss: 7.3669e-04 - 30s/epoch - 3ms/step
Epoch 33/50
10959/10959 - 35s - loss: 3.4655e-04 - val_loss: 7.7535e-04 - 35s/epoch - 3ms/step
Epoch 34/50
10959/10959 - 30s - loss: 3.5163e-04 - val_loss: 6.9589e-04 - 30s/epoch - 3ms/step
Epoch 35/50
10959/10959 - 36s - loss: 3.6533e-04 - val_loss: 7.1176e-04 - 36s/epoch - 3ms/step
Epoch 36/50
10959/10959 - 35s - loss: 3.3477e-04 - val_loss: 6.7198e-04 - 35s/epoch - 3ms/step
Epoch 37/50
10959/10959 - 29s - loss: 3.4120e-04 - val_loss: 6.8567e-04 - 29s/epoch - 3ms/step
Epoch 38/50
10959/10959 - 36s - loss: 3.3053e-04 - val_loss: 7.2224e-04 - 36s/epoch - 3ms/step
Epoch 39/50
10959/10959 - 36s - loss: 3.3720e-04 - val_loss: 7.7018e-04 - 36s/epoch - 3ms/step
Epoch 40/50
10959/10959 - 35s - loss: 3.3583e-04 - val_loss: 8.1989e-04 - 35s/epoch - 3ms/step
Epoch 41/50
10959/10959 - 35s - loss: 3.2631e-04 - val_loss: 7.4394e-04 - 35s/epoch - 3ms/step
Epoch 42/50
10959/10959 - 35s - loss: 3.2230e-04 - val_loss: 8.2110e-04 - 35s/epoch - 3ms/step
Epoch 43/50
10959/10959 - 34s - loss: 3.2935e-04 - val_loss: 8.8571e-04 - 34s/epoch - 3ms/step
Epoch 44/50
10959/10959 - 35s - loss: 3.2474e-04 - val_loss: 6.9232e-04 - 35s/epoch - 3ms/step
Epoch 45/50
10959/10959 - 30s - loss: 3.1958e-04 - val_loss: 6.3976e-04 - 30s/epoch - 3ms/step
Epoch 46/50
10959/10959 - 32s - loss: 3.1371e-04 - val_loss: 7.3848e-04 - 32s/epoch - 3ms/step
Epoch 47/50
10959/10959 - 30s - loss: 3.1649e-04 - val_loss: 7.8478e-04 - 30s/epoch - 3ms/step
Epoch 48/50
10959/10959 - 30s - loss: 3.2053e-04 - val_loss: 7.2122e-04 - 30s/epoch - 3ms/step
Epoch 49/50
10959/10959 - 34s - loss: 3.1754e-04 - val_loss: 6.6753e-04 - 34s/epoch - 3ms/step
Epoch 50/50
10959/10959 - 25s - loss: 3.1484e-04 - val_loss: 7.7200e-04 - 25s/epoch - 3ms/step
```

```
encoder = Model(inputs = input_data_shape,outputs = bottleneck)
encoder.save('encoder.h5')
```

WARNING:tensorflow:Compiled the loaded model, but the compiled metrics have yet to be built. `model.compile_metrics` w

```
from sklearn.metrics import accuracy_score
from xgboost import XGBClassifier
import xgboost as xgb
params = {
    'booster':'gblinear',
    'objective':'binary:logistic',
    'max_depth': 5, 'min_child_weight': 1, 'gamma': 0.1, 'subsample': 0.8,
    'colsample_bytree': 0.8, 'scale_pos_weight': 1,
    'alpha': 10,
    'learning_rate': 0.01,
    'n_estimators':1000,'nthread': 4,'seed': 27
}
xgb_clf = XGBClassifier(**params)
xgb_clf.fit(x1, y1)
y_pred=xgb_clf.predict(x2)
print(y_pred)
accuracy = accuracy_score(y2, y_pred)*100
print("Accuracy before feature selection:-", accuracy)
from sklearn.metrics import f1_score
from sklearn import metrics
print("Confusion Matrix =", metrics.confusion_matrix(y2, y_pred, labels=None,
                                                    sample_weight=None))
print("Recall =", metrics.recall_score(y2, y_pred, labels=None,
```

```
pos_label=1, average='weighted',
sample_weight=None))
print("Precision =", metrics.precision_score(y2, y_pred, labels=None,
pos_label=1, average='weighted',
sample_weight=None))
print("Classification Report =\n", metrics.classification_report(y2, y_pred,
labels=None,
target_names=None,
sample_weight=None,
digits=2,
output_dict=False))

print("F1 Score = ",f1_score(y2, y_pred, average='macro'))
```

```
[1 1 1 ... 1 1 1]
Accuracy before feature selection:- 55.06000097167566
Confusion Matrix = [[ 0 37000]
[ 0 45332]]
Recall = 0.5506000097167566
Precision = 0.30316037070009244
Classification Report =
```

	precision	recall	f1-score	support
0	0.00	0.00	0.00	37000
1	0.55	1.00	0.71	45332
accuracy			0.55	82332
macro avg	0.28	0.50	0.36	82332
weighted avg	0.30	0.55	0.39	82332

```
F1 Score = 0.35508835693695956
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1318: UndefinedMetricWarning: Precision is i
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1318: UndefinedMetricWarning: Precision and
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1318: UndefinedMetricWarning: Precision and
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1318: UndefinedMetricWarning: Precision and
_warn_prf(average, modifier, msg_start, len(result))
```

```
encoder = load_model('encoder.h5')
x1_encode = encoder.predict(x1)
x2_encode = encoder.predict(x2)
params = {
    'booster':'gblinear',
    'objective':'binary:logistic',
    'max_depth': 5, 'min_child_weight': 1, 'gamma': 0.1, 'subsample': 0.8,
    'colsample_bytree': 0.8, 'scale_pos_weight': 1,
    'alpha': 10,
    'learning_rate': 0.01,
    'n_estimators':1000,'nthread': 4,'seed': 27
}
model_final = XGBClassifier(**params)
model_final.fit(x1_encode, y1)
y_pred_new = model_final.predict(x2_encode)
accuracy = accuracy_score(y2, y_pred_new)*100
print("Accuracy after feature selection:-", accuracy)
from sklearn.metrics import f1_score
from sklearn import metrics
print("Confusion Matrix =", metrics.confusion_matrix(y2, y_pred_new, labels=None,
sample_weight=None))
print("Recall =", metrics.recall_score(y2, y_pred_new, labels=None,
pos_label=1, average='weighted',
sample_weight=None))
print("Precision =", metrics.precision_score(y2, y_pred_new, labels=None,
pos_label=1, average='weighted',
sample_weight=None))
print("Classification Report =\n", metrics.classification_report(y2, y_pred_new,
labels=None,
target_names=None,
sample_weight=None,
digits=2,
output_dict=False))

print("F1 Score = ",f1_score(y2, y_pred_new, average='macro'))
```

```
WARNING:tensorflow:No training configuration found in the save file, so the model was *not* compiled. Compile it manua
Accuracy after feature selection:- 71.33799737647574
Confusion Matrix = [[16787 20213]
[ 3385 41947]]
Recall = 0.7133799737647574
Precision = 0.745545159489587
Classification Report =
```

	precision	recall	f1-score	support
0	0.83	0.45	0.59	37000
1	0.67	0.93	0.78	45332
accuracy			0.71	82332
macro avg	0.75	0.69	0.68	82332
weighted avg	0.75	0.71	0.69	82332

F1 Score = 0.683856444205901

